

Introduction-> The Kubernetes Dashboard is a vital tool that simplifies the deployment and management of containerized application. Grafana, a flexible platform for observability, further improves this experience at the same time by providing strong analytics and visualizations. Docker is a platform that facilitates the creation, deployment, and operation of programs through containerization. Code, libraries, and dependencies are all packaged as containers so that they may operate on any machine that supports Docker.

Task 1: Deploy and access the Kubernetes Dashboard and a Web Application Component

Task 1 (A)->Deploy 'Kubernetes Dashboard' on the provided VM with CLI and access/login the Dashboard.

- Dashboard will be available on the following link we need to extract it.
- To extract the link we need to run the command

Command-> kubectl apply -f (and the link)-

<https://raw.githubusercontent.com/kubernetes/dashboard/v2.7.0/aio/deploy/recommended.yaml>

```
student@edge:~/Documents$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.7.0/aio/deploy/recommended.yaml
namespace/kubernetes-dashboard created
serviceaccount/kubernetes-dashboard created
service/kubernetes-dashboard created
secret/kubernetes-dashboard-certs created
secret/kubernetes-dashboard-csrf created
secret/kubernetes-dashboard-key-holder created
configmap/kubernetes-dashboard-settings created
role.rbac.authorization.k8s.io/kubernetes-dashboard created
clusterrole.rbac.authorization.k8s.io/kubernetes-dashboard created
rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard created
clusterrolebinding.rbac.authorization.k8s.io/kubernetes-dashboard created
deployment.apps/kubernetes-dashboard created
service/dashboard-metrics-scraper created
deployment.apps/dashboard-metrics-scraper created
student@edge:~/Documents$ kubectl apply -f Task1Dashboard.yaml
serviceaccount/user created
student@edge:~/Documents$ nano ServiceBinding.yaml
student@edge:~/Documents$ kubectl apply -f ServiceBinding.yaml
clusterrolebinding.rbac.authorization.k8s.io/admin-user created
student@edge:~/Documents$ nano SToken.yaml
student@edge:~/Documents$ kubectl apply -f SToken.yaml
secret/admin-user created
student@edge:~/Documents$ kubectl get secret admin-user -n kubernetes-dashboard
-o jsonpath={".data.token"} | base64 -d
```

Fig 1.1

- Task 1(B)-> Creating a Service Account named admin-user in the Kubernetes-dashboard namespace involves using Kubernetes configuration file.
- Running command with kubectl tool and can be done using a .yaml file to define the service account.

Command to create a .yaml file->[nano dashboard.yaml\(filename\).yaml](#) figure 1.2 shows the same

Code Snippet->

```
student@edge:~$ cat dashborad.yaml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: admin-user
  namespace: kubernetes-dashboard
```

Fig 1.2

- **Task 1(C)->**Managing access in Kubernetes dashboard, with default role like cluster -admin are often preconfigured in the cluster.
- This role has some set of permission that provide to access.
- ClusterRoleBinding need to be created to grant some sort of permission to service account manually if the service account does not have the permissions.

Command to create ClusterRolebinding file manually [nano dashboard1\(filename\).yaml](#)

Code Snippet->

```
student@edge:~$ cat dashboard1.yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: admin-user
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: admin-user
  namespace: kubernetes-dashboard
```

Fig 1.3

- Figure 1.3 shows the code to create cluster role binding
- Managing access controls in Kubernetes, certain default roles like **cluster-admin** are often preconfigured in the cluster. These roles come with a set of permissions that provide extensive access across the entire cluster.
- **Task 1(D)->**To access the Kubernetes dashboard we need to access the create a token so for this there are two ways.
 - (i) Temporary token using command-> [kubectl -n kubernetes-dashboard create token admin-user](#) one time access the token next time again we need to generate new token to access the dashboard.
 - (ii) Create a token with the secret which bound the service account and the token will be saved in the Secret
- After token is created we need to enable the port 8081 on which the Kubernetes dashboard will be accessible using this command->[kubectl proxy](#) .

```
student@edge:~$ kubectl get secret admin-user -n kubernetes-dashboard -o jsonpath='{.data.token}' | base64 -d
eyJhbGciOiJSUzIiNisImlpZCI6I1BPNXZwUHFVfMGZ5TTJaekIid24Tkc3M21TVGY5b2l2Vm9iaHo3LXRv0UiFQ.eyJpc3MiOiJrdJlc5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXJuZXRLcy5pb9ZJ2awNLYNbj3VudC9uYm3bhY2U0iJrdJlc5ldGVzL2FyZCIsImt1YmVybmr0ZXMuaW8vc2VydmljZWfjY291bnQvc2VjcmV0Lm5hbWU0iJhZGpbil1c2VylwiLa3ViZXJuZXRLcy5pb9ZJ2awNLYNbj3VudC9zZXJaWNllWFjY291bnQubmFtZSI6ImFkbwlulXVzXIlClJrdJlc5ldGVzLm1vL3NlcnZpY2VhY2Nvd50L3NLcnZpY2UtYNNjb3VudC51aWQ10iIzN2E3MjQwZS0zNzkyLTQyMWMtYmQzYy1jMGEwZjMyNDfkNjIiLCJzdWii0iJzeXN0Zw06c2VydmljZWfjY291bn06a3ViZXJuZXRLcy1YXNm0Ym9hcm06YWRtaW4tdXNlcj9. VxgKItiy_nwSibtYL9A5Fnj5SC00hdjydt6R83FALss9m7m01dxRMU1w4Eo1LWws1vA5K_rmFTDD5YoMEH6QYp_1FO15Rt63Z2tGEmwCNfgBabcPhuU04C8dxkuwlLlufB1010_Ard1SFTjS07AgC014fwdJ9NBm2R8AIT3LU3CwEhepJVatR6FggyAOajvgPC02aj0ltlpxTQA1LFBZ0bFFsaI_P2-h7Dymigmxn-Fyax_lAIgyqNQ5xK5rs3aNPlqza002_ps1gT7RhIpOKd1YUcJmyxLT4v72drI0B1v3s6IdvBMRer7cvA_-6if1bJxt@student@edge
:$
student@edge:~$ kubectl proxy
Starting to serve on 127.0.0.1:8001
```

Fig 1.4

- Figure 1.4 shows how token is generated and 8081 port in enables to access the dashboard .
- Figure 1.5 shows we need to provide the token

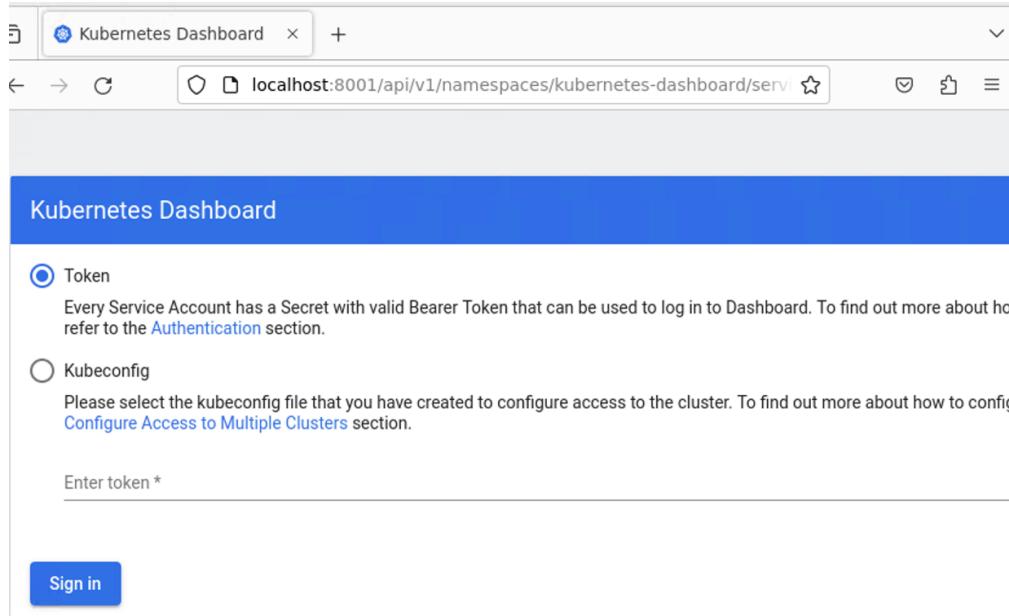


Fig 1.4

- **Task 1(E)->**Deploy an instance of the Docker image “nclcloudcomputing/javabenchmark” via CLI.
- We need to deploy it manually through command line using .yaml file
- Figure 1.5 shows the code to pull the image.

Command ->[nano javabenchmark1.yaml](#) to pull the docker image.

Command->[kubectl apply -f javabenchmark1.yaml](#).

Code Snippet->

```
student@edge:~$ cat javabenchmark1.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javabenchmarkapp-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: javabenchmarkapp
  template:
    metadata:
      labels:
        app: javabenchmarkapp
    spec:
      containers:
        - name: javabenchmarkapp-container
          image: nclcloudcomputing/javabenchmarkapp
          ports:
            - containerPort: 8080
```

Fig 1.5

- **Task 1(F)->**Deploy a NodePort service so that the web app is accessible via <http://localhost:30000/primecheck>.
- NordPort is used so that we can access it locally localhost:30000/primecheck
Command->`kubectl apply -f javabenchmarkapp.yaml`

```
student@edge:~$ cat javabenchmarkapp.yaml
apiVersion: v1
kind: Service
metadata:
  name: javabenchmarkapp-service
spec:
  type: NodePort
  selector:
    app: javabenchmarkapp
  ports:
    - name: http
      protocol: TCP
      port: 8080
      targetPort: 8080
      nodePort: 30000
```

Fig 1.6

- Figure 1.7 primecheck is accessible on 30000 port

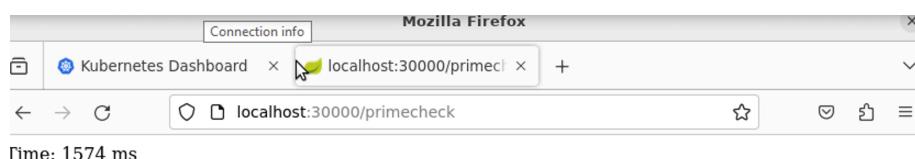


Fig 1.7

Task 2 Deploy the monitoring stack of Kubernetes.

- **Task 2(A)->**Enable observability services from micro8s addon.
- To enable services command is microk8s enable observability.
- Figure 2.2 and 2.3 shows cases to enable observability.
- **Task 2(B)->**Grafana Services has been edited to type NordPort so that we can access it locally and on port 30001 figure 2.4.
- And Grafana can also be access using the cluster IP.
- Figure 2.5 shows all the resources status and running on which port/Cluster IP.
- **Task 2(C)->** Grafana login through **localhost:30001**

```
Infer repository core for addon observability
Addon core/dns is already enabled Time: 1574 ms
Addon core/helm3 is already enabled
Addon core/hostpath-storage is already enabled
Enabling observability
Release "kube-prom-stack" does not exist. Installing it now.
NAME: kube-prom-stack
LAST DEPLOYED: Mon Nov 27 16:06:44 2023
NAMESPACE: observability
STATUS: deployed
REVISION: 1
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace observability get pods -l "release=kube-prom-stack"

Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager and Prometheus instances using the Operator.
Release "loki" does not exist. Installing it now.
NAME: loki
LAST DEPLOYED: Mon Nov 27 16:07:19 2023
NAMESPACE: observability
STATUS: deployed
REVISION: 1
NOTES:
The Loki stack has been deployed to your cluster. Loki can now be added as a datasource in Grafana.

See http://docs.grafana.org/features/datasources/loki/ for more detail.
Release "tempo" does not exist. Installing it now.
NAME: tempo
LAST DEPLOYED: Mon Nov 27 16:07:24 2023
NAMESPACE: observability
STATUS: deployed
REVISION: 1
TEST SUITE: None
[sudo] password for student:

Note: the observability stack is setup to monitor only the current nodes of the MicroK8s cluster.
For any nodes joining the cluster at a later stage this addon will need to be set up again.

Observability has been enabled (user/pass: admin/prom-operator)
```

Fig 2.1

```
student@edge:~$ kubectl edit service kube-prom-stack-grafana --namespace observability
```

Fig 2.2

```

! Please edit the object below. Lines beginning with a '#' will be ignored,
! and an empty file will abort the edit. If an error occurs while saving this file will be
! reopened with the relevant failures.

apiVersion: v1
kind: Service
metadata:
  annotations: {}
    meta.helm.sh/release-name: kube-prom-stack
    meta.helm.sh/release-namespace: observability
  creationTimestamp: "2023-11-23T11:45:05Z"
  labels:
    - except: requests.exceptions.Timeout
  app.kubernetes.io/instance: kube-prom-stack
  app.kubernetes.io/managed-by: Helm
  app.kubernetes.io/name: grafana
  app.kubernetes.io/version: 9.3.8
  helm.sh/chart: grafana-6.51.2
  name: kube-prom-stack-grafana
  namespace: observability
  resourceVersion: "81630"
  uid: 3106af3a-ef4c-4174-aa44-0378f4307ef9
spec:
  clusterIP: 10.152.183.195
  clusterIPs:
  - 10.152.183.195
  sessionAffinity: None
  selector:
    app.kubernetes.io/instance: kube-prom-stack
  ports:
  - name: http-web
    targetPort: 8080
    port: 80
    protocol: TCP
    # or = LoadGenerator(target, frequency)
    # targetPort: 3000 generate load()
  selector:
    app.kubernetes.io/instance: kube-prom-stack

```

Fig 2.3

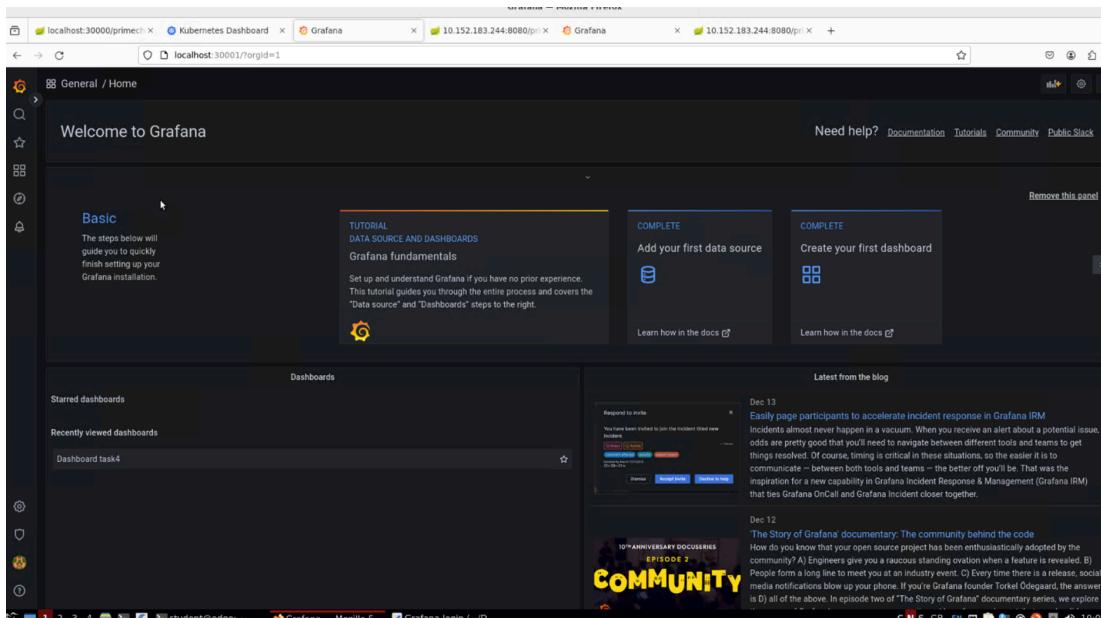


Fig 2.4

```

student@edge:~$ kubectl get all --namespace=observability
NAME                                         READY   STATUS    RESTARTS   AGE
pod/kube-prom-stack-prometheus-node-exporter-h524d   1/1     Running   1 (14d ago)  28d
pod/alertmanager-kube-prom-stack-kube-prome-alertmanager-0   2/2     Running   2 (14d ago)  28d
pod/kube-prom-stack-kube-state-metrics-6c586bf4c8-7qvfd   1/1     Running   1 (14d ago)  28d
pod/kube-prom-stack-kube-prome-operator-64ffd55b77-dlggv   1/1     Running   1 (14d ago)  28d
pod/loki-promtail-x67r8   1/1     Running   1 (14d ago)  28d
pod/tempo-9   2/2     Running   2 (14d ago)  28d
pod/prometheus-kube-prom-stack-kube-prome-prometheus-0   2/2     Running   2 (14d ago)  28d
pod/kube-prom-stack-grafana-6c47f548d6-95zwx   3/3     Running   3 (14d ago)  28d
pod/loki-0   1/1     Running   1 (14d ago)  28d

NAME          TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)
AGE
service/kube-prom-stack-prometheus-node-exporter ClusterIP  10.152.183.171 <none>           9100/TCP
20d
The steps below will help you understand what each service does and how it interacts with other components in the stack.
service/kube-prom-stack-kube-prome-operator ClusterIP  10.152.183.137 <none>           443/TCP
20d
Grafana fundamentals
service/kube-prom-stack-kube-prome-alertmanager ClusterIP  10.152.183.37 <none>           9093/TCP
20d
Alerting and alerting rules
service/kube-prom-stack-kube-prome-prometheus ClusterIP  10.152.183.131 <none>           9090/TCP
20d
Metrics and monitoring
service/kube-prom-stack-kube-state-metrics ClusterIP  10.152.183.107 <none>           8080/TCP
20d
Metrics and monitoring
service/alertmanager-operated ClusterIP  None            <none>           9093/TCP, 9094/TCP, 9094/UDP
20d
Learn how in the docs
service/prometheus-operated ClusterIP  None            <none>           9090/TCP
20d
Metrics and monitoring
service/loki-headless ClusterIP  None            <none>           3100/TCP
20d
Logs and log processing
service/loki-memberlist ClusterIP  None            <none>           7946/TCP
20d
Logs and log processing
service/loki ClusterIP  10.152.183.204 <none>           3100/TCP
20d
Logs and log processing
service/tempo ClusterIP  10.152.183.170 <none>           3100/TCP, 16687/TCP, 16686/TCP, 6831/UDP, 6832/UDP, 14268/TCP, 14
TCP, 9411/TCP, 55680/TCP, 55681/TCP, 4317/TCP, 4318/TCP, 55678/TCP
20d
Metrics and monitoring
service/kube-prom-stack-grafana NodePort   10.152.183.195 <none>           80:30001/TCP
20d

NAME          DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
daemonset.apps/kube-prom-stack-prometheus-node-exporter 1         1         1         1         1         1         <none>           20d
daemonset.apps/loki-promtail 1         1         1         1         1         1         <none>           20d

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/kube-prom-stack-kube-prome-operator 1/1     1           1           20d
deployment.apps/kube-prom-stack-kube-state-metrics 1/1     1           1           20d
deployment.apps/kube-prom-stack-grafana   1/1     1           1           20d

NAME          DESIRED   CURRENT   READY   AGE
replicaset.apps/kube-prom-stack-kube-prome-operator-64ffd55b77 1         1         1         20d

```

Fig 2.5

Task 3 Load Generator

- **Task 3(A)**->Write a load generator with following specification.
 - (a) Accepts two configurable values either via a config file or environment variables. target (The address for the load generation) and frequency (Request per second)
 - (b) Generate web request to the target at the specified frequency
 - (c) Collect 2 types of metrics. Average response time and accumulated number of failures
 - (d) Request should timeout if it takes more than 10 seconds. Counted as failures
 - (e) Test results need to be printed to the console
 - (f) There are no requirements in programming language

Create a python file using Command-> [nano loadgenerator\(filename\).py](#)

Code Snippet->

```
student@edge:~$ cat loadgenerator.py
import requests
import time
import os
Network
class LoadGenerator:
    def __init__(self, target, frequency):
        self.target = target
        self.frequency = frequency
        self.total_requests = 0
        self.total_failures = 0
        self.total_response_time = 0

    def generate_load(self):
        while True:
            start_time = time.time()
            try:
                response = requests.get(self.target, timeout=10)
                if response.status_code != 200:
                    self.total_failures += 1
            except requests.exceptions.Timeout:
                self.total_failures += 1

            end_time = time.time()
            self.total_requests += 1
            self.total_response_time += end_time - start_time

            time.sleep(1 / self.frequency)

            if self.total_requests % 10 == 0: # Print metrics every 10 requests
                self.print_metrics()

    def print_metrics(self):
        avg_response_time = self.total_response_time / self.total_requests if self.total_requests > 0 else 0
        print(f"Metrics after {self.total_requests} requests:")
        print(f"Average Response Time: {avg_response_time:.4f} seconds")
        print(f"Accumulated Number of Failures: {self.total_failures}\n")

if __name__ == "__main__":
    target = os.getenv("TARGET_ADDRESS", "http://10.152.183.244:8080/primecheck") # Replace with default target
    frequency = float(os.getenv("REQUESTS_PER_SECOND", 10)) # Replace with default frequency

    load_generator = LoadGenerator(target, frequency)
    load_generator.generate_load()
```

Fig 3.1

- Figure 3.2 shows the response after running the **loadgenerator** python file manually through command line using command -> [kubectl apply -f loadgenerator.py\(filename\)](#) .

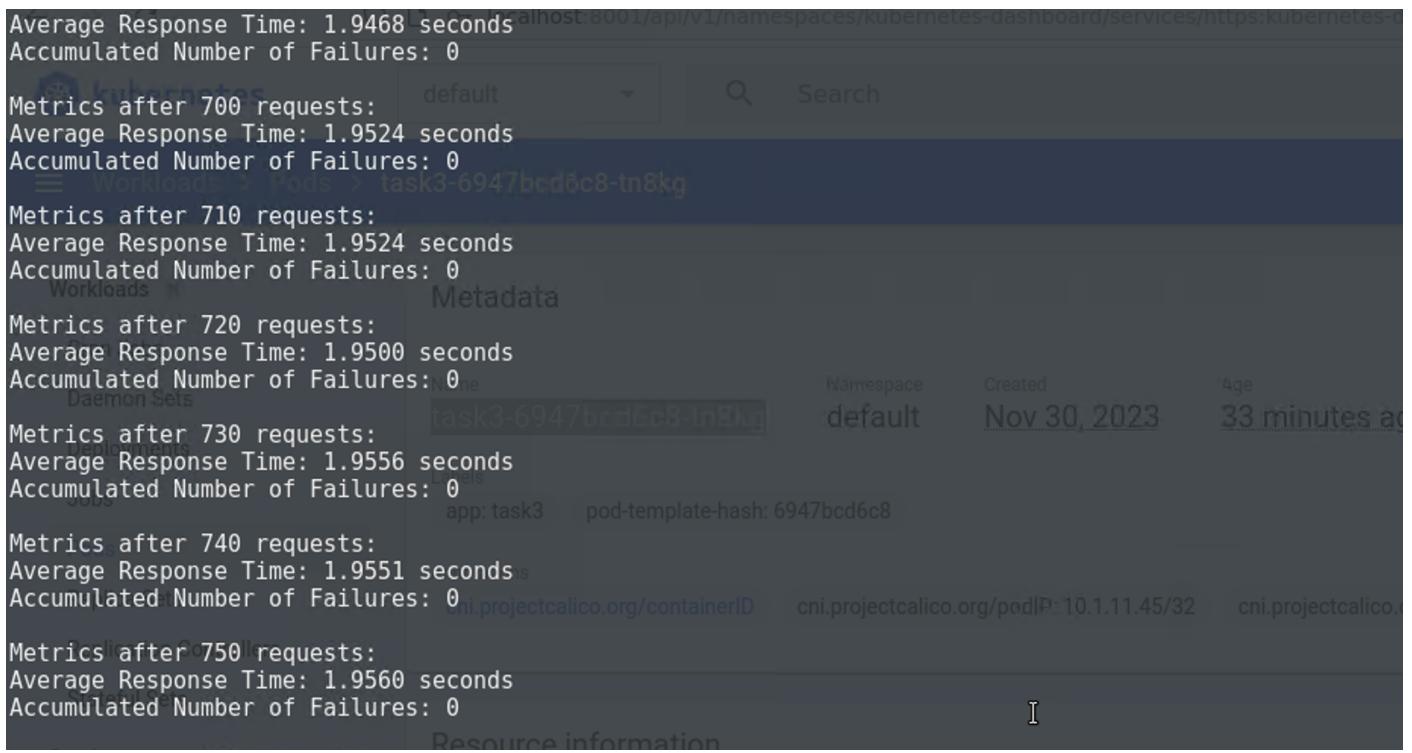


Fig 3.2

- **Task 3(B)->**Once the **loadgenerator** python file is done, next step is to create a standalone docker image of the same with name load-generator.

Command->**docker build -t load-generator .**



Fig 3.3

- Figure 3.3 shows docker image is created with name load-generator.
- Docker image **load-generator** need to pushed on **local-registry on 32000 port**.

Command ->**docker run -d -p 32000:5000 --restart=always --name registry registry:latest**

- After that we need to tag the image.

Command->**docker tag load-generator local:32000/load-generator**

- Push image to local repository

Command->**docker push localhost:32000/load-generator**

```
student@edge:~$ dockerfile run -d -p 32000:5000 --restart=always --name registry registry:2
dockerfile: command not found
student@edge:~$ docker run -d -p 32000:5000 --restart=always --name registry registry:2
Unable to find image 'registry:2' locally
2: Pulling from library/registry
96526aa774ef: Pull complete
834bccaa730c: Pull complete
87a69098c0a9: Pull complete
afc17120a9f7: Pull complete
e5ac04f3acf5: Pull complete
Digest: sha256:8a60daaa55ab0df4607c4d8625b96b97b06fd2e6ca8528275472963c4ae8afa0
Status: Downloaded newer image for registry:2
60864dd28e93daf3f296329a310e7fd509f417b7fb8898e65e243995e490769b
```

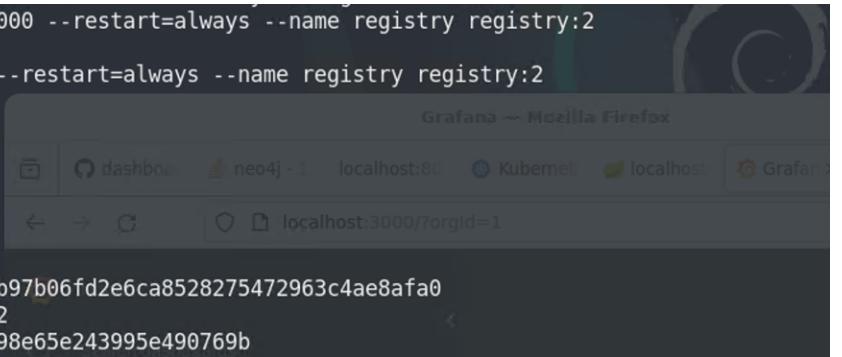
A screenshot of a Mozilla Firefox browser window. The title bar says "Grafana — Mozilla Firefox". The address bar shows "localhost:3000?orgId=1". The main content area displays a complex Grafana dashboard with multiple panels, graphs, and data series, typical of a monitoring or analytics interface.

Fig 3.4

- Figure 3.4 docker image **load-generator** is pushed to local registry on 32000

Task 4 Monitor Benchmarking Results

- **Task 4(A)->**Deploy load-generator service created in task 3 through command line.
- First .yaml file need to be created using command ->[nano task3.yaml\(filename\)](#) .

Code Snippet->

```
student@edge:~$ cat task3.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: task3
spec:
  replicas: 1
  selector:
    matchLabels:
      app: task3
  template:
    metadata:
      labels:
        app: task3
    spec:
      containers:
        - name: task3
          image: localhost:32000/load-generator
          ports:
            - containerPort: 8080
          env:
            - name: target
              value: "http://10.152.183.244:8080/primecheck"
            - name: frequency
              value: "10.0"
            command: ["python", "loadgenerator.py"]

student@edge:~$ kubectl apply -f task3.yaml
deployment.apps/task3 configured
```

Fig 4.1

- Figure 4.1 shows that service has been deployed with image localhost:32000/load-generator
- **Task 4(B)->** Create a new dashboard on Grafana and add 2 new panels which should contain queries of CPU/memory usage of the web application
- First panel with ->CPU usage.
- Figure 4.3 shows the memory panel graph

Metric ->container_memory_usuage_bytes

Label Filter ->container-> javabenchmarkapp-conatiner .

- Memory used in bytes on Y-axis and time on X-axis.
- Figure 4.4 shows the CPU panel graph with

Metric ->container_cpu_usuage_seconds_tools

Label Filter ->container -> javabenchmarkapp-conatiner .

- Figure 4.2 shows the both the panel and from that we can understand the memory and CPU usage.
- Figure 4.6 shows the logs



Fig 4.2



Fig 4.3

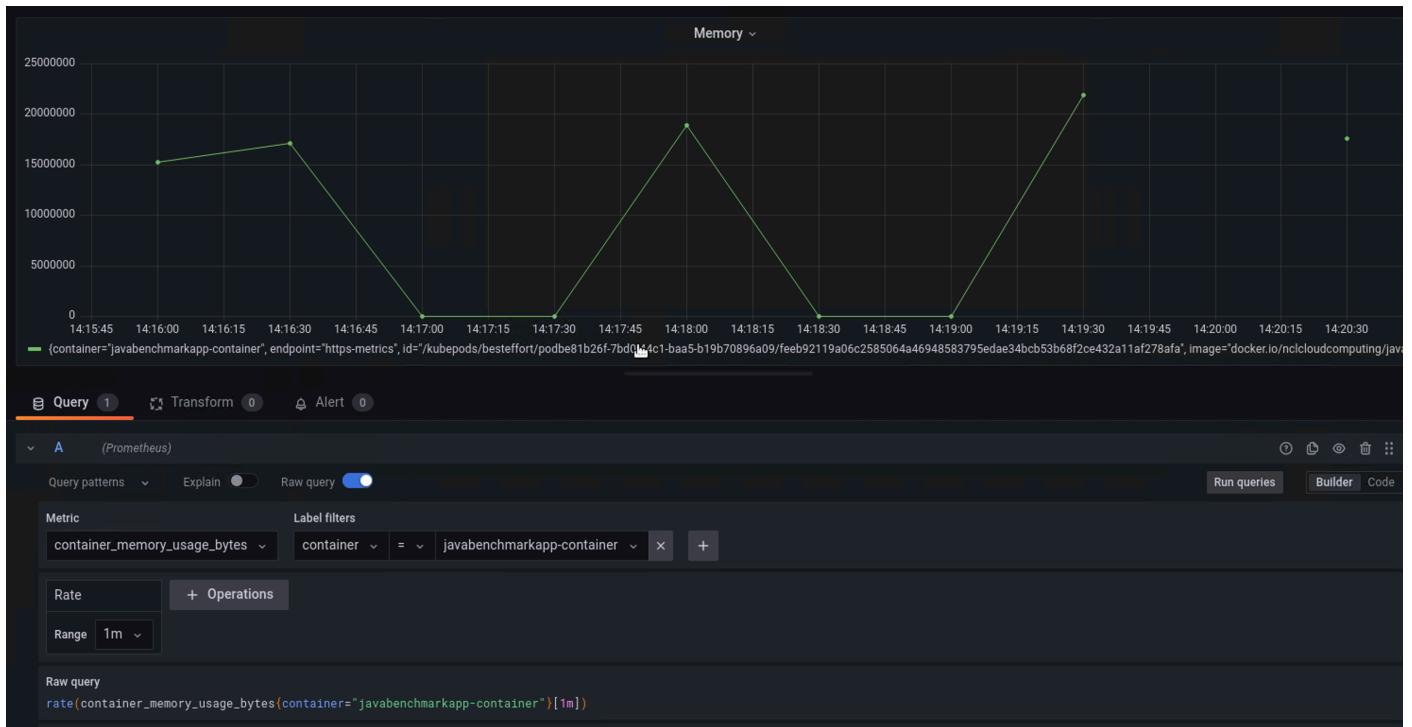


Fig 4.4

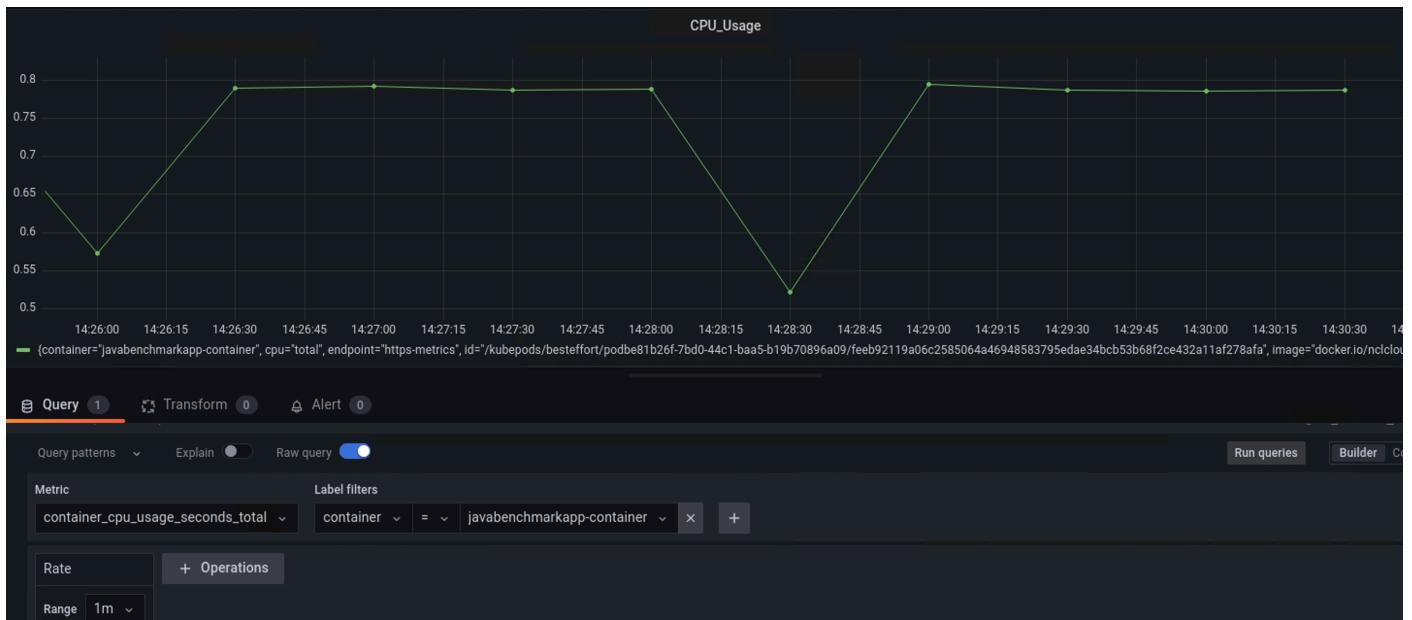


Fig 4.5

Logs from task3 in task3-6947bcd...

```
Metrics after 570 requests:  
Average Response Time: 1.9096 seconds  
Accumulated Number of Failures: 0  
Metrics after 580 requests:  
Average Response Time: 1.9141 seconds  
Accumulated Number of Failures: 0  
Metrics after 590 requests:  
Average Response Time: 1.9158 seconds  
Accumulated Number of Failures: 0  
Metrics after 600 requests:  
Average Response Time: 1.9207 seconds  
Accumulated Number of Failures: 0  
Metrics after 610 requests:  
Average Response Time: 1.9220 seconds  
Accumulated Number of Failures: 0  
Metrics after 620 requests:  
Average Response Time: 1.9213 seconds  
Accumulated Number of Failures: 0  
Metrics after 630 requests:  
Average Response Time: 1.9234 seconds  
Accumulated Number of Failures: 0  
Metrics after 640 requests:  
Average Response Time: 1.9263 seconds  
Accumulated Number of Failures: 0  
Metrics after 650 requests:  
Average Response Time: 1.9295 seconds  
Accumulated Number of Failures: 0  
Metrics after 660 requests:  
Average Response Time: 1.9344 seconds  
Accumulated Number of Failures: 0  
Metrics after 670 requests:
```

Fig 4.6

Conclusion-> For controlling Kubernetes clusters, Kubernetes Dashboard offers an intuitive user interface. Docker uses containerization to make application deployment easier while maintaining consistency between environments. Grafana provides robust system metrics monitoring and visualization features that improve understanding of application performance and infrastructure health. When combined, they improve efficiency and dependability in contemporary software development and operations by streamlining management, deployment, and monitoring.